

Module 4 — Subagents

Lesson 4.3

Prompting subagents effectively

Brief a subagent like a smart colleague who just walked into the room. Include everything they need; cap what you expect back.

Why it matters

Subagents have no memory of your conversation. The brief is the whole world. The same task with a good vs. bad brief can be the difference between a useful report and a wasted minute.

The anatomy of a good brief

A brief should answer four questions:

- 1 **What you're trying to accomplish.** Not just the immediate task — the *why*.
- 2 **What you've already learned or ruled out.** Don't make the agent repeat work.
- 3 **What constraints apply.** Tools available, files in scope, format expected.
- 4 **What you want back.** Form, length, level of detail.

Template

Task: <one-sentence goal>

Context:

- <why this matters / what surrounds it>
- <what's already been tried or ruled out>

Scope:

- <files / dirs / log paths in scope>
- <files explicitly out of scope, if relevant>

Constraints:

- <max time / max files to read / read-only / etc.>

Deliverable:

- <exact format you want>
- <word/length cap>
- <fields/sections required>

You don't need every section every time, but the elements matter.

Examples

Bad brief

"Find the bug in the user auth code."

The agent has no idea what bug, where the code is, what's been tried, or what to return. Expect a meandering report.

Better brief

Task: Find the source of a 401 error returned to legit users after token refresh.

Context:

- Users get 401 on `/api/me` ~1 in 10 requests, only after token refresh.
- The refresh itself succeeds (we see the new token in the cookie).
- I've ruled out clock skew (system time is within 100ms of NTP).
- I suspect a race between the refresh endpoint and downstream calls.

Scope:

- `src/auth/` (the auth code)
- `src/middleware/auth.ts` (the middleware that checks tokens)
- Logs at `/var/log/api.log` for the last hour

Constraints:

- Read-only. Don't change anything.
- Don't read files outside `src/auth/` and `src/middleware/`.

Deliverable:

- Most likely cause in one sentence.
- Top 3 file:line locations to investigate, with one-sentence why each.
- A concrete next experiment (test to write, log to grep, query to run).
- Under 200 words.

This brief lets the agent be useful.

Patterns that work

"Report in under N words"

Forces the agent to prioritize. Without a cap, agents tend to produce 1500-word reports of which 300 are signal.

"Cite file:line for every claim"

Forces grounding in code, not inference. A claim without a citation is a guess; you can hold the agent to that standard.

"Tell me what would change your mind"

For research and review tasks, ask the agent to explicitly note what evidence would flip its conclusion. Surfaces low-confidence claims.

"Don't speculate"

For diagnostic work. Without it, the agent will offer "this could also be caused by..." trailers that pad the report and bury the lead.

"Self-contained reply"

Reminds the agent its reply will be read by someone who didn't see its thinking. No "as I mentioned above" — there is no above for the reader.

Anti-patterns

"Based on the previous agent's findings..."

The new agent has no findings to refer to. Always include the relevant findings in the new brief.

Vague output specs

"Give me a thorough analysis" → expect a long meandering reply. "Give me a 5-line verdict with file:line citations" → expect a useful one.

Conflating exploration with implementation

If you want both research and a fix, brief two agents (or one with general-purpose and explicit tool grants). Mixing them in a vague brief produces neither.

Asking the agent to ask you questions

Subagents can't loop with you mid-task. If your brief has unresolved ambiguity, resolve it before spawning. If absolutely necessary, tell the agent "if X is ambiguous, assume Y and proceed."

Length and time budgeting

A brief shouldn't be longer than a page. If you find yourself writing more than that, either:

- The task is too complex — split it into multiple briefs
- You're under-trusting the agent — trust the system prompt for general behavior

Cap word counts in the deliverable, not in the brief.

Try it

- 1 Pick a research task you'd normally do inline.
- 2 Write a brief using the template. Include all four sections.
- 3 Spawn an Explore subagent with it. Note the quality.
- 4 Re-do with just "find X." Compare.

The difference will be stark. Internalize the pattern.

Gotchas

- **Briefs are not the place for the team's coding conventions.** Those belong in CLAUDE.md or the agent definition.
- **Don't include irrelevant context "just in case."** It dilutes signal and burns the agent's window.
- **Be explicit about format.** "Markdown table" beats "structured output." "JSON with these keys" beats "JSON."
- **Subagents won't push back on weird requests.** If you brief them to do something dumb, they'll attempt it. The check is on you, before spawning.

Further reading

- 4.4 Parallel orchestration — how briefs differ for parallel fan-out
- 2.6 When to delegate — when to even spawn

Next

→ 4.4 Parallel orchestration